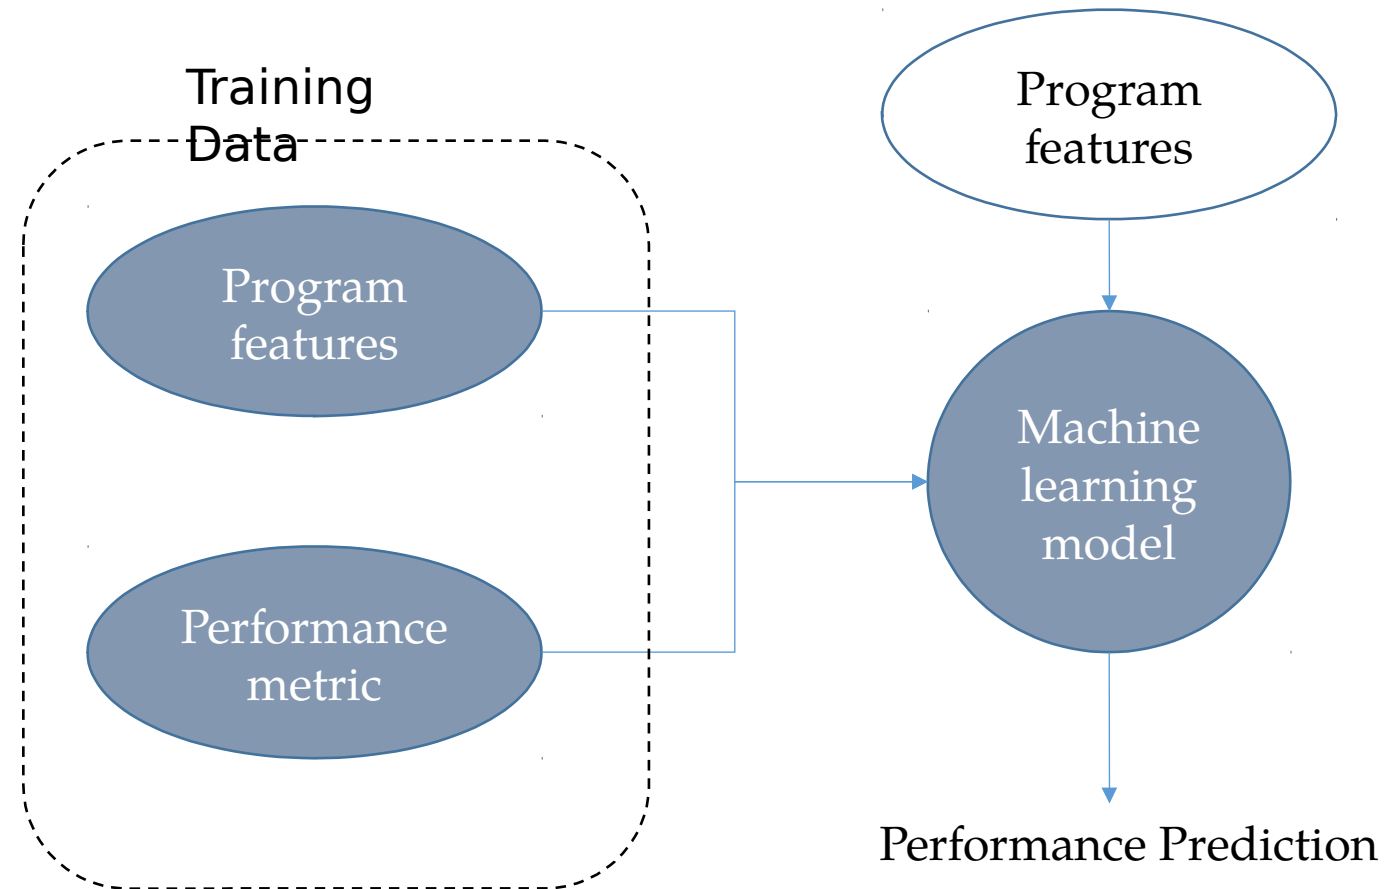


# Performance Bug Diagnosis using Performance Predictions

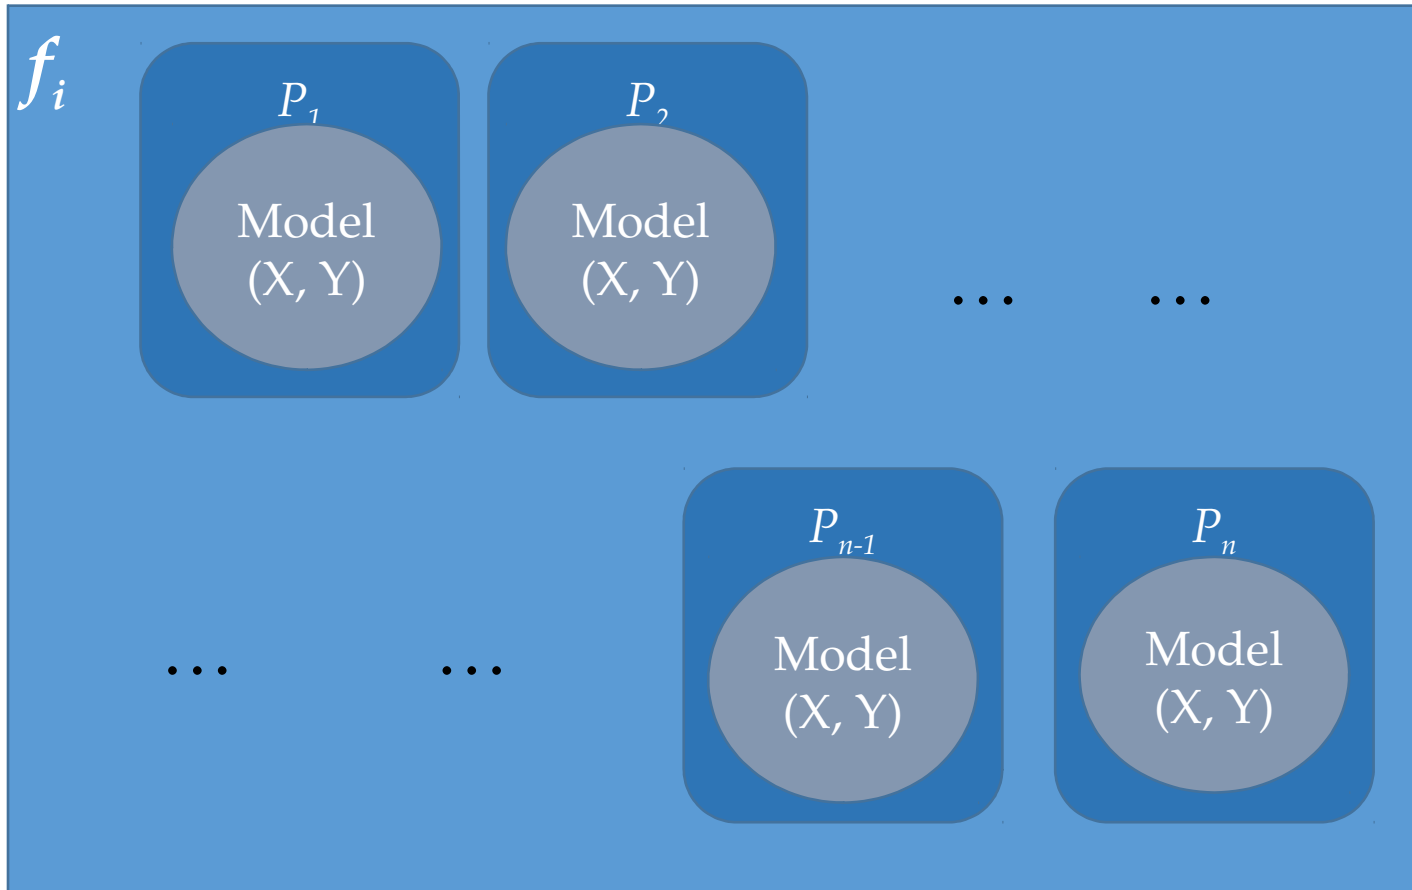
# Program Performance Prediction



# Program Performance Prediction

- Training data for different performance bugs
  - Simulating a buggy execution
  - Collect program features and corresponding performance metric
- Construct performance prediction model for specific performance bug scenario
  - Machine learning / Statistical models

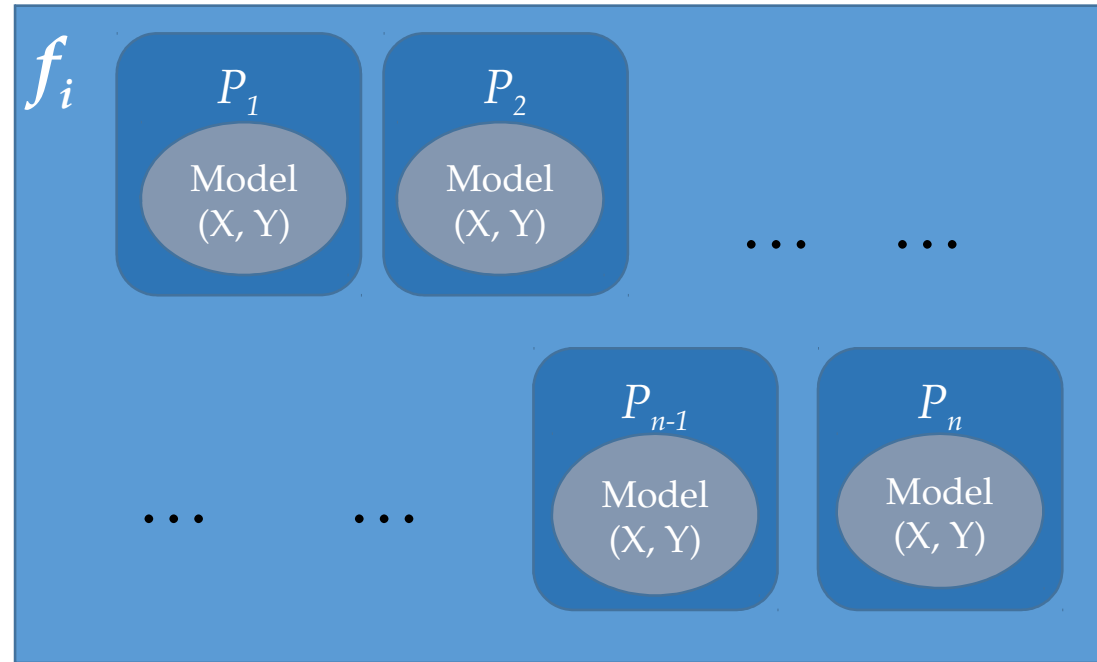
# Modeling Performance Bug



- $F = \{ f_1, f_2 \dots \}$ 
  - Set of functions
- $X = \langle x_1, x_2, \dots \rangle$ 
  - Set of features
- $Y$ 
  - Performance metric
- $P = \{ p_1, p_2, \dots \}$ 
  - Performance bugs

# Performance Bug Diagnosis

- Predict performance for current features,
  - $X = \langle x_1, x_2, \dots \rangle$
- Performance predictions for all bug models



Prediction from bug models of  $P_i ==$  Actual observed performance  Bug  $P_i$


# Modeling Performance Bug

- Set of functions,  $F = \{ f_1, f_2 \dots \}$ 
  - Hot functions
- Set of features,  $X = \langle x_1, x_2, \dots \rangle$ 
  - Code features : CPI, load/store and branch count etc.
  - Data features : Loop count, branch miss rate, cache miss rate etc.
- Performance metric,  $y$ 
  - Execution time
- Performance bugs  $P = \{ p_1, p_2, \dots \}$ 
  - High cache miss rate or branch misprediction rate, high cache or lock contentions etc.

# Challenges

- What are the fixed set of features?
- How to generate training data for specific performance bug?
- How to create a performance prediction model accurate enough for performance anomaly detection?
- What is an effective granularity of program to predict performance?

# Modeling Lock Contention Problem




- Execution time of critical section (CS):

$$T_{cs} = t_2 - t_1$$

- $\Delta t = 0$   lock contention
- $T_{cs} > t_3 - t_1$   lock contention



# Modeling Lock Contention Problem



- Profile  $T_{cs}$  for a contention free execution
  - Positive training example
- Estimate  $T_{cs}$  for contention scenerio
  - Negative training example
  - $T_{cs} = t_3 - t_1$



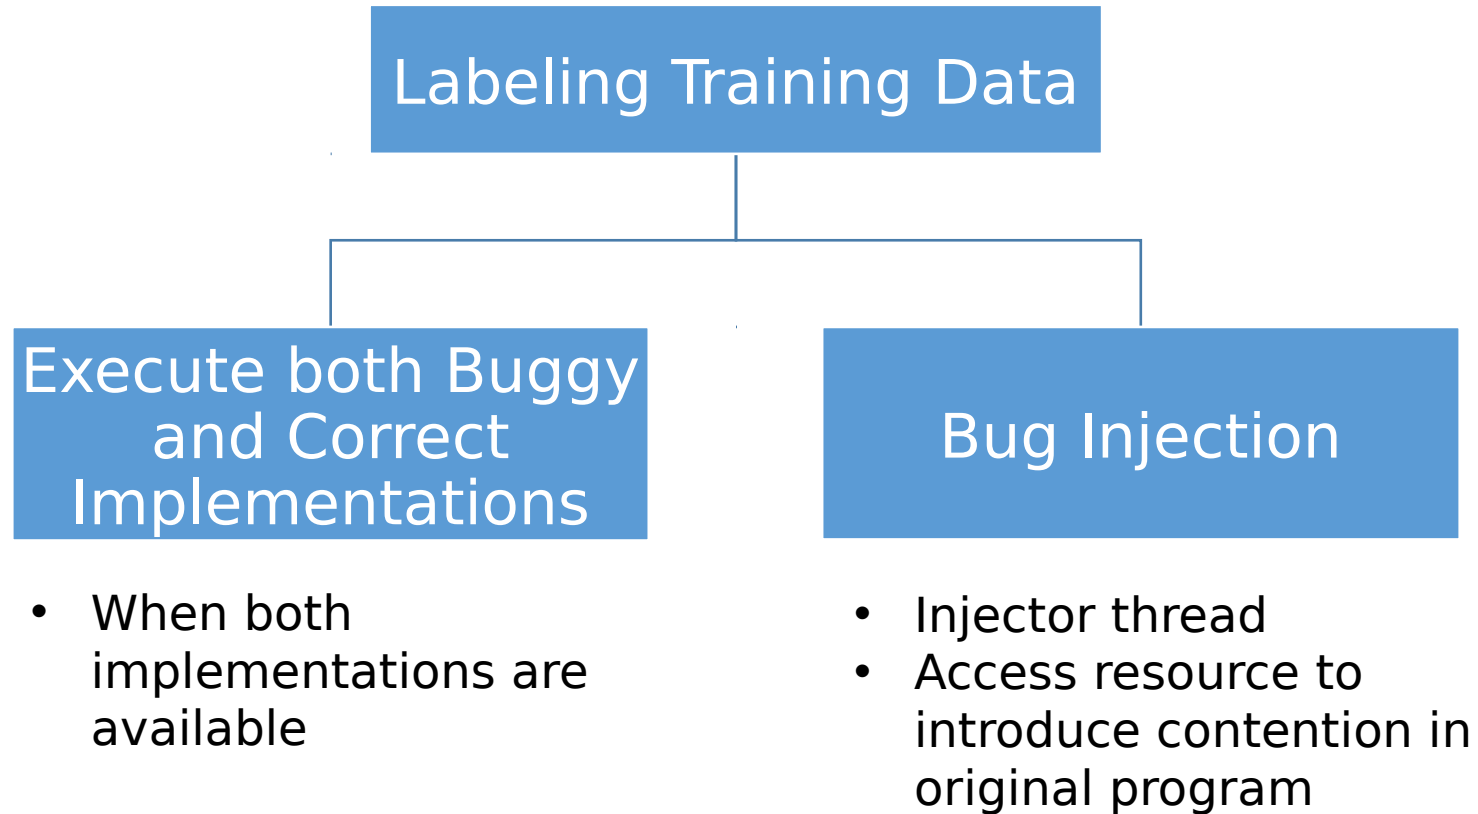
# Performance Anomaly Detection

- Collect features for a specific function(or block of code)
- Features for Cache Contention Problem:
  - Cache invalidation rate
  - Number of threads
  - IPC

# Machine Learning Model

- Training Data :
  - Features collected from dynamic instances of the function(or block of code)
  - Execute the program with
    - Different input sizes
    - Different thread numbers
    - Different machine configurations

# Machine Learning Model



**Use a threshold difference in IPC to label data as an instance of anomalous execution**

# Experiments

- PARSEC & Phoenix benchmarks
  - Different input sizes (small, large)
  - Different thread numbers (8, 16 & 32 threads)
  - Different machine configurations ( 8, 16 & 32 core)
- Machine learning models: SVM, MLP & Decision tree

# Experiments

- Labeling is done by running both correct & buggy implementations
- Anomaly due to cache contention detected(known issues):
  - Streamcluster – false sharing in function *double pgain(...)*
  - Dataset size: 26850
  - Cross validation scores
    - SVM : ~99%
    - MLP : ~99 %
    - Decision Tree: ~99%